



# EE 046746 - Technion - Computer Vision

Dahlia Urbach

## Tutorial 09 - Introduction to 3D Deep Learning

---

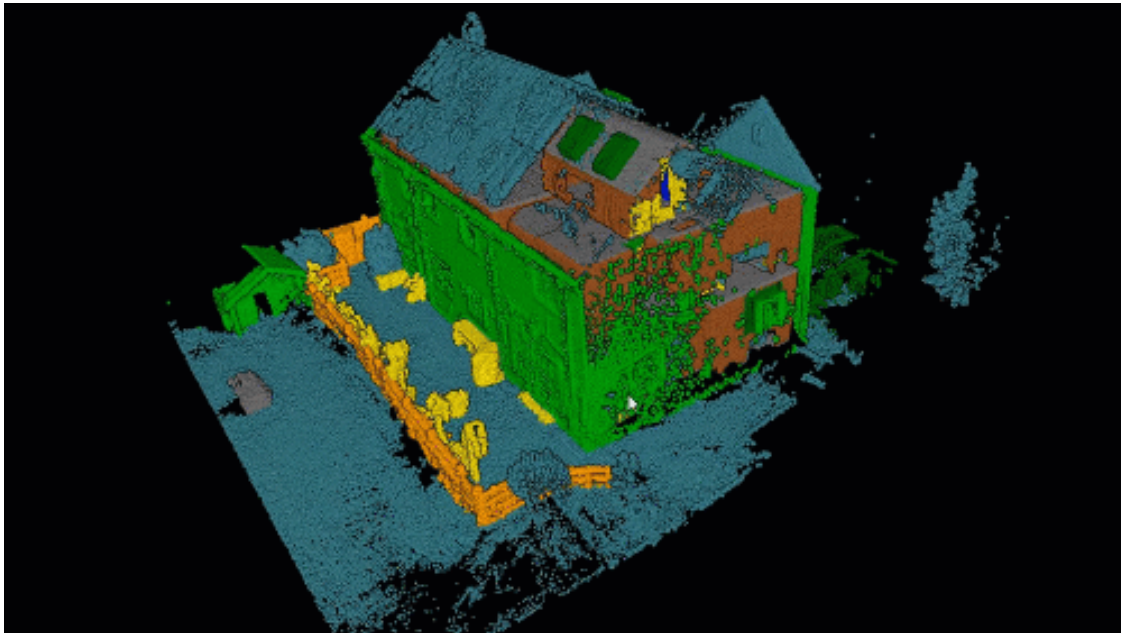


Image source (<https://towardsdatascience.com/the-future-of-3d-point-clouds-a-new-perspective-125b35b558b9>).

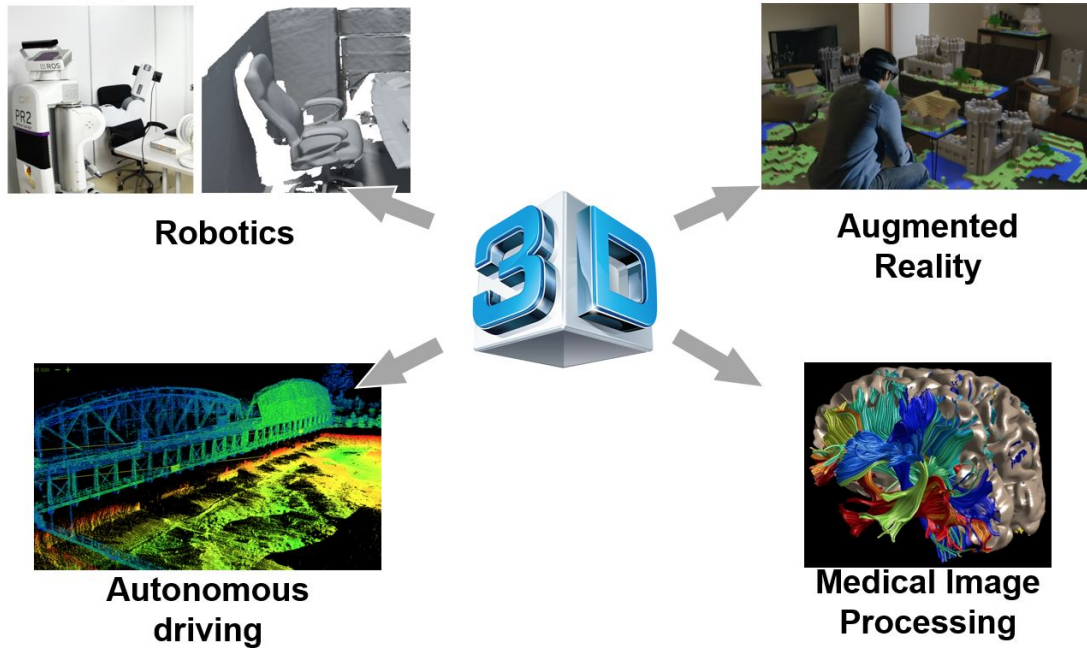
LiDAR point cloud support | Feature Highlight | Unreal Engine (<https://www.youtube.com/watch?v=R-ZXdAEGbiw&feature=youtu.be>).



## Agenda

---

- [Depth Cameras - Quick overview](#)
  - Stereo Cameras - Next Week
  - [Time of Flight](#)
- [3D Deep Learning](#)
  - [Voxels](#)
  - [Multi-View](#)
  - [Point Clouds](#)
- [3D Applications](#)
- [Recommended Tools](#)
- [Recommended Videos](#)
- [Credits](#)



```
In [8]: # imports for the tutorial
import numpy as np
import matplotlib.pyplot as plt
import time

# pytorch imports
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset, ConcatDataset
# import torchvision
```



## Depth Cameras

- Stereo Cameras - Next week
- Time of flight

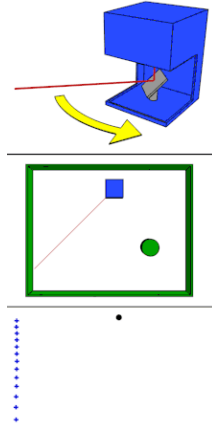


## Time of Flight Cameras

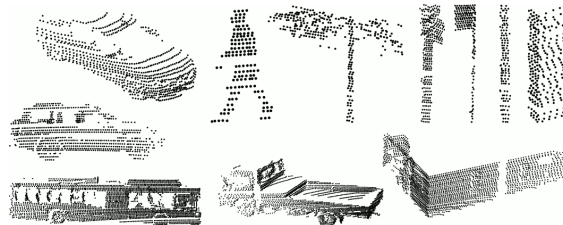
- Light travels at approximately a constant speed  $c = 3 \times 10^8$  (meters per second).
- Measuring the time it takes for light to travel over a distance once can infer distance.
- Can be categorized into two types:
  1. Direct TOF - switch laser on and off rapidly.
  2. Indirect TOF - send out modulated light, then measure phase difference to infer depth.

## 1. Direct - TOF

- **Light Detection And Ranging** (LiDAR) probably best example in computer vision and robotics.
- High-energy light pulses limit influence of background illumination.
- However, difficulty to generate short light pulses with fast rise and fall times.
- High-accuracy time measurement required.
- Prone to motion blur.
- Sparser as objects grow in distance.

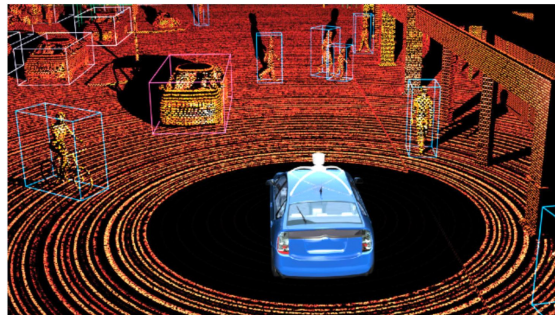


Gif source - Wikipedia (<https://en.wikipedia.org/wiki/Lidar>).

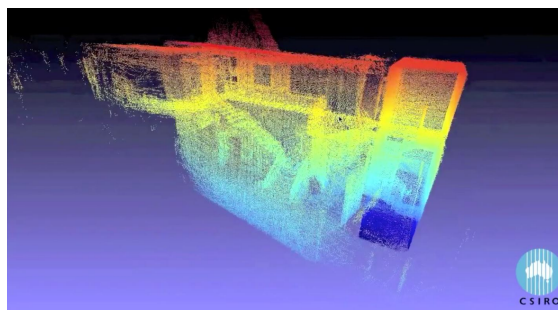


Sydney Dataset (<http://www.acfr.usyd.edu.au/papers/SydneyUrbanObjectsDataset.shtml>).

## Autonomous Car - LiDAR



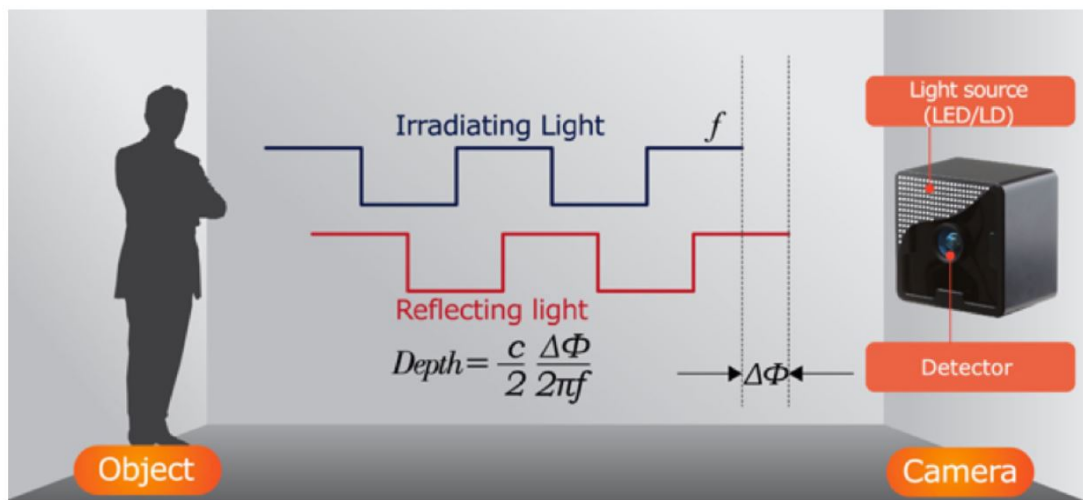
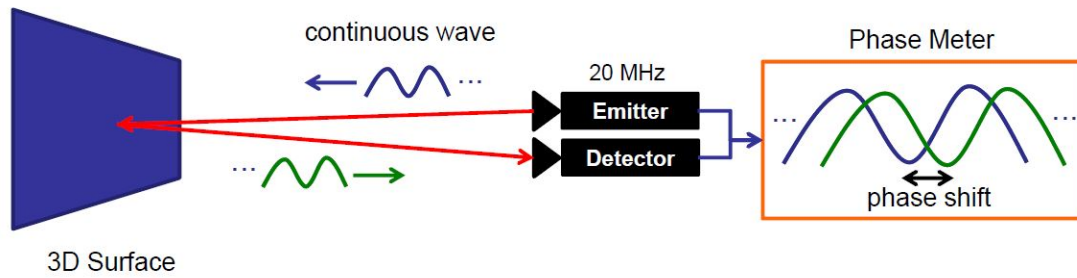
## SLAM + LIDAR - Zebedee

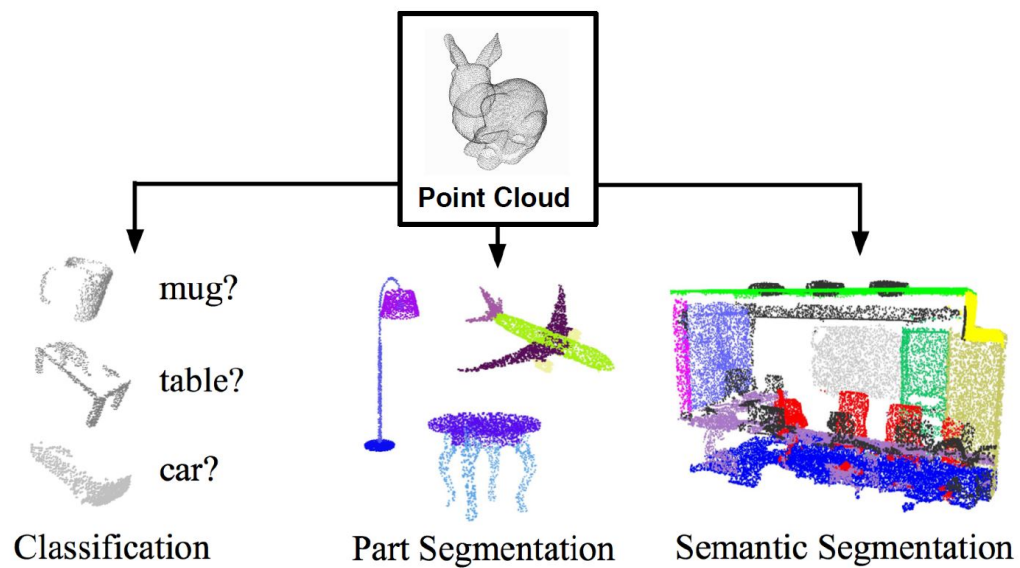


zebedee (<https://research.csiro.au/robotics/zebedee/>).

## 2. Indirect - TOF

- Continuous light waves instead of short light pulses.
- Modulation in terms of frequency of sinusoidal waves.
- Detected wave after reflection has shifted phase.
- Phase shift proportional to distance from reflecting surface.





- Classification
- Semantic segmentation
- Part segmentation
  - Each point belongs to a specific part of the object
- ...

Qi, Charles R., et al. "Pointnet: Deep learning on point sets for 3d classification and segmentation." CVPR. 2017.



### Questions

- What are the differences between 2D image and a point cloud?
- Why it might be hard to feed a point cloud as neural network input?
- What are the benefits of using a point cloud?



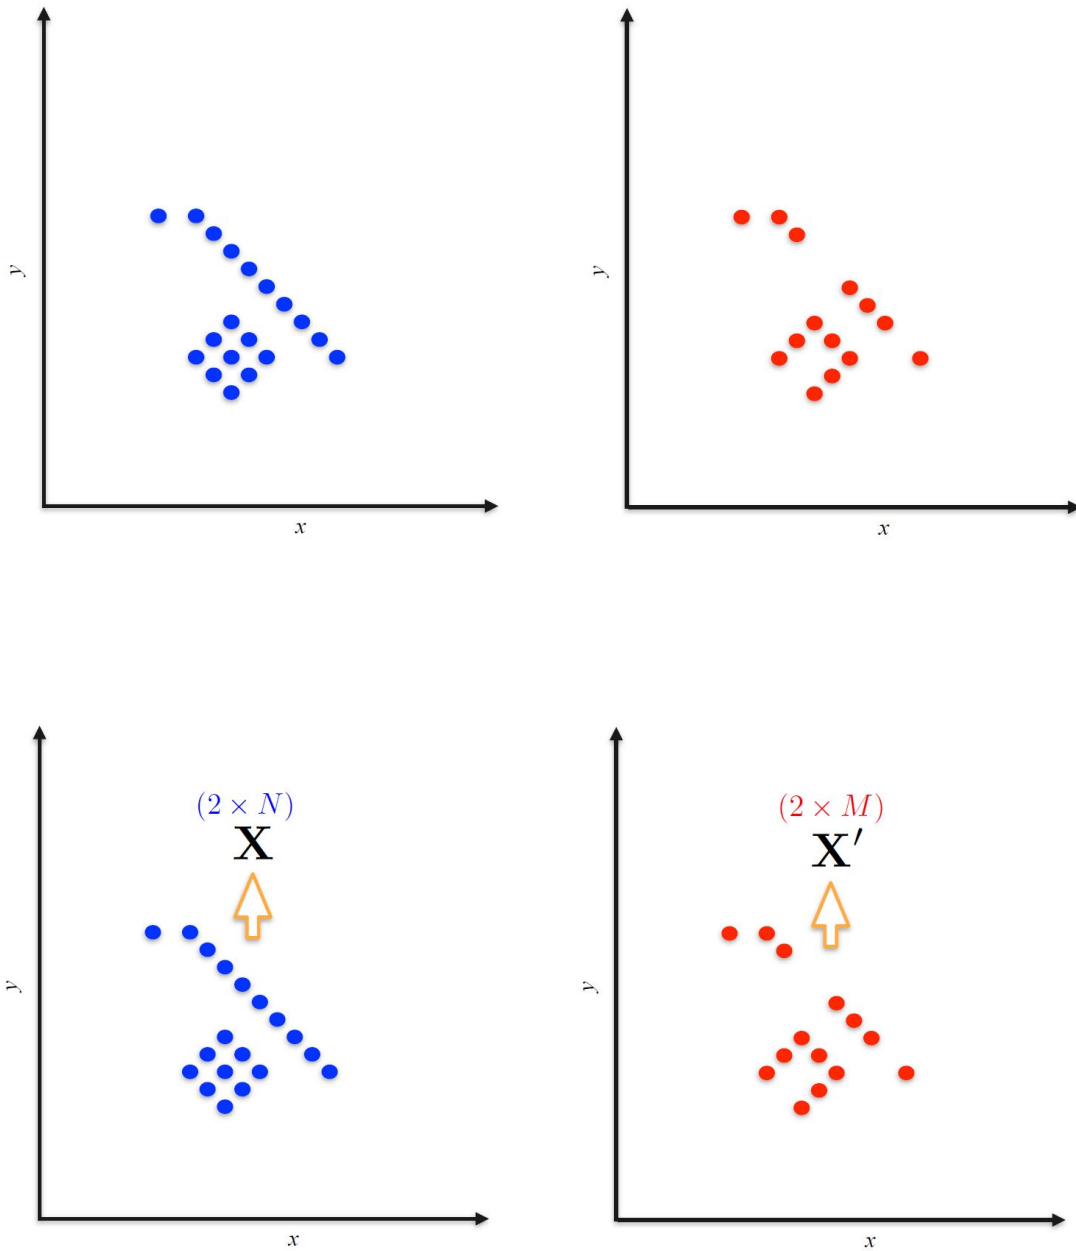
### Point Clouds Problems

- Point Clouds Vary in Size (not constant)
- Unordered Input
  - Data is unstructured (no grid)
  - Data is invariant to point ordering (permutations)

### Other Point Clouds Challenges

- Missing data
- Noise
- Rotations

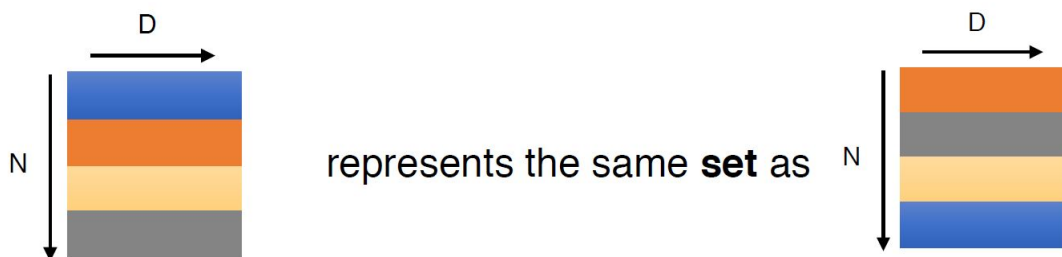
### Problem - Point Clouds Vary in Size



- Different point clouds represent the same object

### Problem - Unordered Input

Point cloud:  $N$  **orderless** points, each represented by a  $D$  dim vector



How many semi-equal representations?

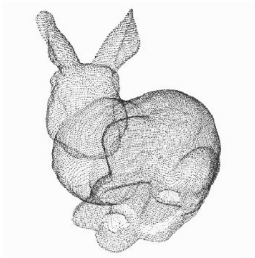
Model needs to be invariant to  $N!$  permutations



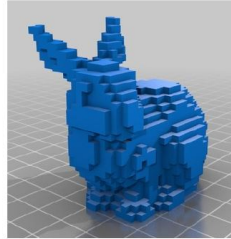
## Alternate 3D Representations - Representations

Solution:

- Convert the raw point clouds into Voxels or multiple 2D RGB(D) images



Point Cloud



Volumetric

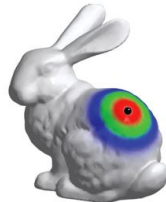


Projected View  
RGB(D)

Another 3D representation (not in this course):



Part Assembly



Mesh (Graph CNN)

$$F(x) = 0$$

Implicit Shape



## Voxelization

Idea: generalize 2D convolutions to regular 3D grids

- The straightforward approach: transform the point clouds into a voxel grid by rasterizing and use 3D CNNs



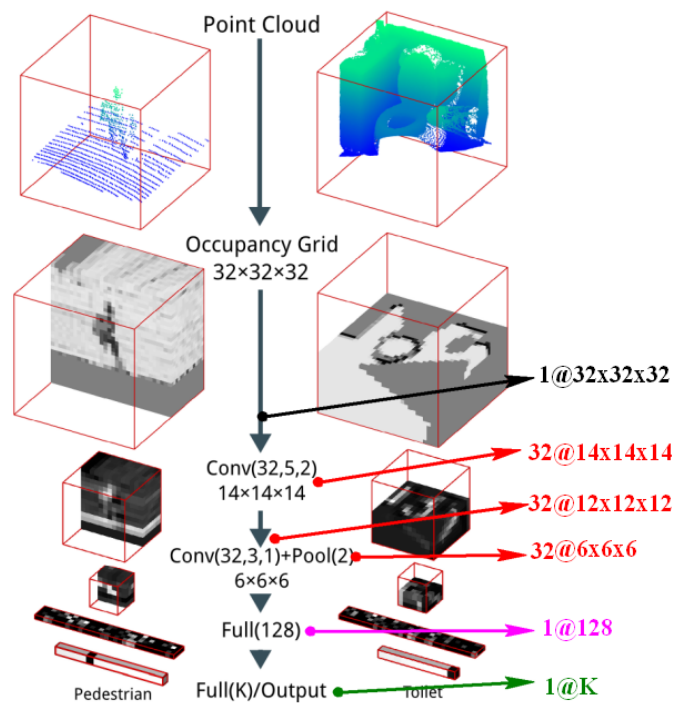
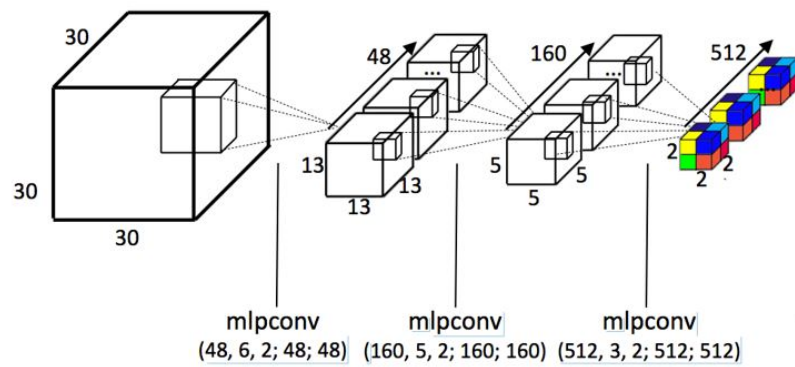
Voxel grid is a 3D grid of equal size volumes (voxels), can be occupied by:

- Binary 0/1 - Is there any point within the voxel?
- Weighted - The amount of point located within each voxel

Usually we use binary occupancy



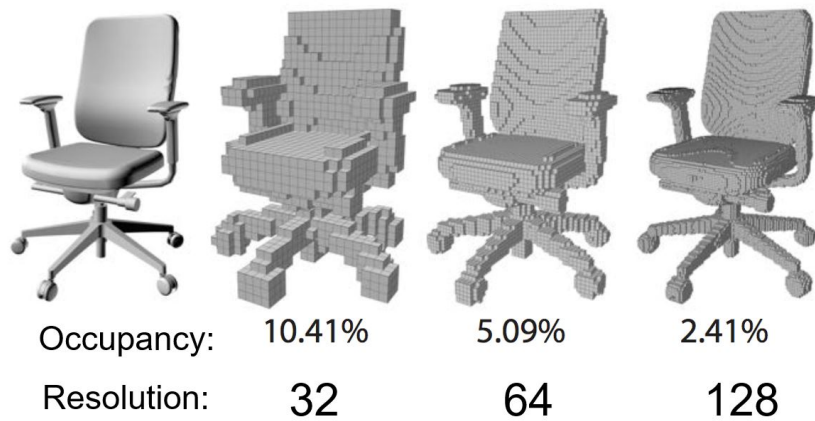
3D convolution uses 4D kernels



Maturana, Daniel, and Sebastian Scherer. Voxelnet: A 3d convolutional neural network for real-time object recognition. IROS, 2015.



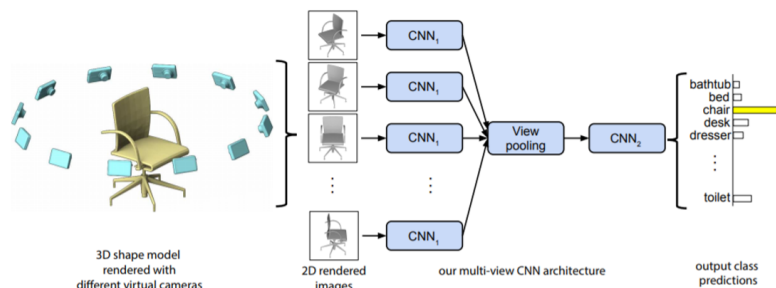
- Large memory cost
- Slow processing time
- Limited spatial resolution
- Quantization artifacts



## Multi-View Approach

Idea: Transform the problem into a well known domain (3D→2D)

- The multi-view approach: project multiple views to 2D and use CNN to process
  - How many views do we need? (Another hyper parameter)



CNN<sub>1</sub> - We can use pre-trained networks to extract features followed by fine tune layers

H. Su, S. Maji, E. Kalogerakis, and E. Learned-Miller. Multiviewconvolutional neural networks for 3d shape recognition. CVPR, 2015.

## Apply Deep Learning Directly on 3D Point Clouds

Idea: Most of the raw 3D data are point clouds - Solve the problems!

Note: Point Clouds Problems:

- Point Clouds Vary in Size (not constant)
- Unordered Input
  - Data is unstructured (no grid)
  - Data is invariant to point ordering (permutations)



### PointNet

**Permutation Invariance: Symetric Function**

$$f(x_1, x_2, \dots, x_n) \equiv f(x_{\pi_1}, x_{\pi_2}, \dots, x_{\pi_n}), x_i \in R^D$$

$\pi$  is a different permutation

Example:

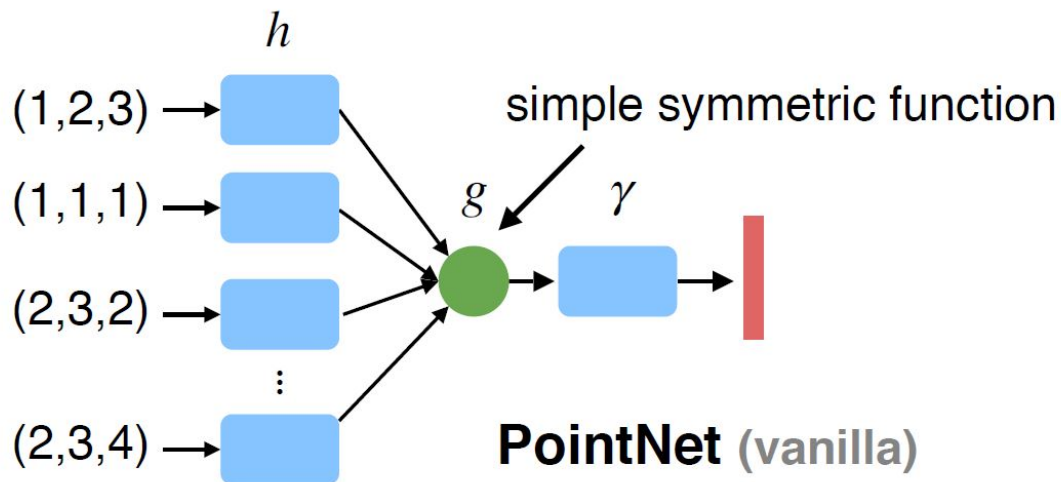
$$f(x_1, x_2, \dots, x_n) = \max\{x_1, x_2, \dots, x_n\}$$
$$f(x_1, x_2, \dots, x_n) = x_1 + x_2 + \dots + x_n$$

- How can we construct a family of symmetric functions by neural networks?

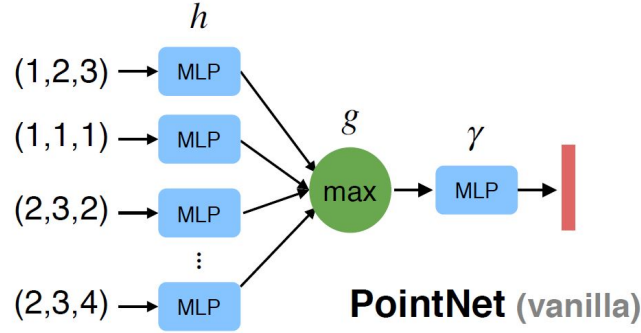
Observe:

$$f(x_1, x_2, \dots, x_n) = \gamma \circ g(h(x_1), \dots, h(x_n))$$

is symmetric if  $g$  is symmetric



Empirically, we use multi-layer perceptron (MLP) and max pooling:



Input MLP:

$$h(x_i) : R^3 \rightarrow R^D$$

We can look at it as  $D$  functions  $\{h_k\}_{k=1}^D$  operate on each point where

$$h_k(x_i) : R^3 \rightarrow R^1$$

Pooling layer:

$$g(h(x_1), \dots, h(x_n)) : R^{N \times D} \rightarrow R^D$$

We apply the pooling over all points for each function  $h_k$ .

$$g(h_k(x_1), \dots, h_k(x_n)) : R^{N \times 1} \rightarrow R^1$$

Classification MLP:

$$\gamma \circ g(h(x_1), \dots, h(x_n)) : R^D \rightarrow R^{D_{\text{NumClasses}}}$$

Shared MLP implementation "trick":

- Use conv layers : Number of filters  $C_{out}$ , each filter size is  $1 \times C_{in}$ .
- Input:  $R^{N \times C_{in}}$
- Output:  $R^{N \times C_{out}}$

MLP:

$$h : R^{C_{in}} \rightarrow R^{C_1} \rightarrow \dots \rightarrow R^{C_{out}}$$

```

In [7]: class Tnet(nn.Module):
    def __init__(self, k=3):
        super().__init__()
        self.k=k
        self.conv1 = nn.Conv1d(k,64,1)
        self.conv2 = nn.Conv1d(64,128,1)
        self.conv3 = nn.Conv1d(128,1024,1)
        self.fc1 = nn.Linear(1024,512)
        self.fc2 = nn.Linear(512,256)
        self.fc3 = nn.Linear(256,k*k)

        self.bn1 = nn.BatchNorm1d(64)
        self.bn2 = nn.BatchNorm1d(128)
        self.bn3 = nn.BatchNorm1d(1024)
        self.bn4 = nn.BatchNorm1d(512)
        self.bn5 = nn.BatchNorm1d(256)

    def forward(self, input):
        # input.shape == (bs,n,3)
        bs = input.size(0)
        xb = F.relu(self.bn1(self.conv1(input)))
        xb = F.relu(self.bn2(self.conv2(xb)))
        xb = F.relu(self.bn3(self.conv3(xb)))
        pool = nn.MaxPool1d(xb.size(-1))(xb)
        flat = nn.Flatten(1)(pool)
        xb = F.relu(self.bn4(self.fc1(flat)))
        xb = F.relu(self.bn5(self.fc2(xb)))

        #initialize as identity
        init = torch.eye(self.k, requires_grad=True).repeat(bs,1,1)
        if xb.is_cuda:
            init=init.cuda()
        matrix = self.fc3(xb).view(-1,self.k,self.k) + init
        return matrix

class Transform(nn.Module):
    def __init__(self):
        super().__init__()
        self.input_transform = Tnet(k=3)
        self.feature_transform = Tnet(k=64)
        self.conv1 = nn.Conv1d(3,64,1)

        self.conv2 = nn.Conv1d(64,128,1)
        self.conv3 = nn.Conv1d(128,1024,1)

        self.bn1 = nn.BatchNorm1d(64)
        self.bn2 = nn.BatchNorm1d(128)
        self.bn3 = nn.BatchNorm1d(1024)

    def forward(self, input):
        matrix3x3 = self.input_transform(input)
        # batch matrix multiplication
        xb = torch.bmm(torch.transpose(input,1,2), matrix3x3).transpose(1,2)

        xb = F.relu(self.bn1(self.conv1(xb)))

        matrix64x64 = self.feature_transform(xb)
        xb = torch.bmm(torch.transpose(xb,1,2), matrix64x64).transpose(1,2)

        xb = F.relu(self.bn2(self.conv2(xb)))
        xb = self.bn3(self.conv3(xb))
        xb = nn.MaxPool1d(xb.size(-1))(xb)
        output = nn.Flatten(1)(xb)
        return output, matrix3x3, matrix64x64

class PointNet(nn.Module):
    def __init__(self, classes = 10):
        super().__init__()
        self.transform = Transform()
        self.fc1 = nn.Linear(1024, 512)
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, classes)

        self.bn1 = nn.BatchNorm1d(512)
        self.bn2 = nn.BatchNorm1d(256)
        self.dropout = nn.Dropout(p=0.3)
        self.logsoftmax = nn.LogSoftmax(dim=1)

    def forward(self, input):
        xb, matrix3x3, matrix64x64 = self.transform(input)
        xb = F.relu(self.bn1(self.fc1(xb)))
        xb = F.relu(self.bn2(self.dropout(self.fc2(xb))))

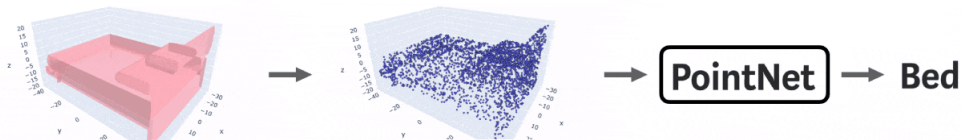
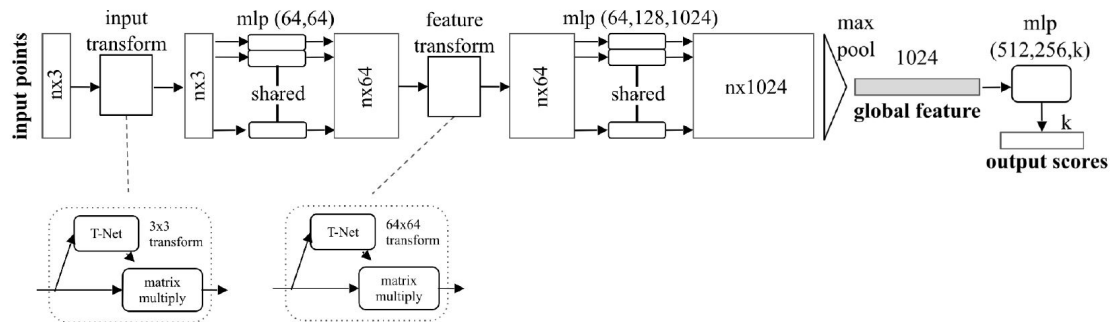
```

```
output = self.fc3(xb)
return self.logsoftmax(output), matrix3x3, matrix64x64
```

[Code source \(https://github.com/nikitakaraevv/pointnet/blob/master/nbs/PointNetClass.ipynb\)](https://github.com/nikitakaraevv/pointnet/blob/master/nbs/PointNetClass.ipynb) - Full implementation of PointNet Classification (can be opened in Colab)

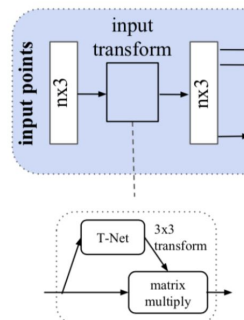
[Deep Learning on Point clouds: Implementing PointNet in Google Colab \(https://towardsdatascience.com/deep-learning-on-point-clouds-implementing-pointnet-in-google-colab-1fd65cd3a263\)](https://towardsdatascience.com/deep-learning-on-point-clouds-implementing-pointnet-in-google-colab-1fd65cd3a263) - Nikita Karaev

## PointNet Classification Network



[Image source \(https://towardsdatascience.com/deep-learning-on-point-clouds-implementing-pointnet-in-google-colab-1fd65cd3a263\)](https://towardsdatascience.com/deep-learning-on-point-clouds-implementing-pointnet-in-google-colab-1fd65cd3a263)

## Transformation Invariance



Learn transformation matrix to improve task performance. We want our network to be invariante to rigid transformation of the object.

- Practically, more augmentation over the training dataset also solved the transformation invariance

[Image source \(https://medium.com/@luis\\_gonzales/an-in-depth-look-at-pointnet-111d7efd4a1a\)](https://medium.com/@luis_gonzales/an-in-depth-look-at-pointnet-111d7efd4a1a)

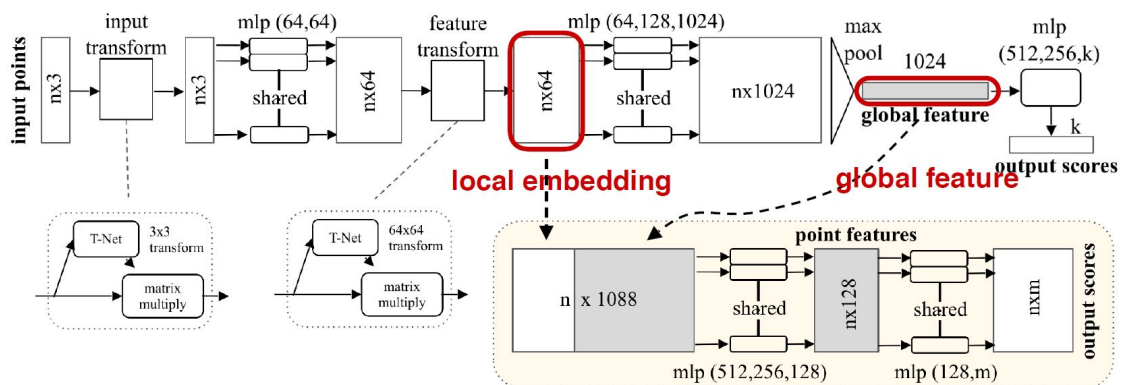
Qi, Charles R., et al. "PointNet: Deep learning on point sets for 3d classification and segmentation." CVPR 2017.

## Results - Classification

	input	#views	accuracy avg. class	accuracy overall
SPH [12]	mesh	-	68.2	
3D CNNs	3DShapeNets [29]	volume	77.3	84.7
	VoxNet [18]	volume	83.0	85.9
	Subvolume [19]	volume	86.0	89.2
	LFD [29]	image	75.5	-
	MVCNN [24]	image	90.1	-
	Ours baseline	point	72.6	77.4
	Ours PointNet	point	86.2	89.2

Qi, Charles R., et al. "PointNet: Deep learning on point sets for 3d classification and segmentation." CVPR 2017.

## PointNet Segmentation Network

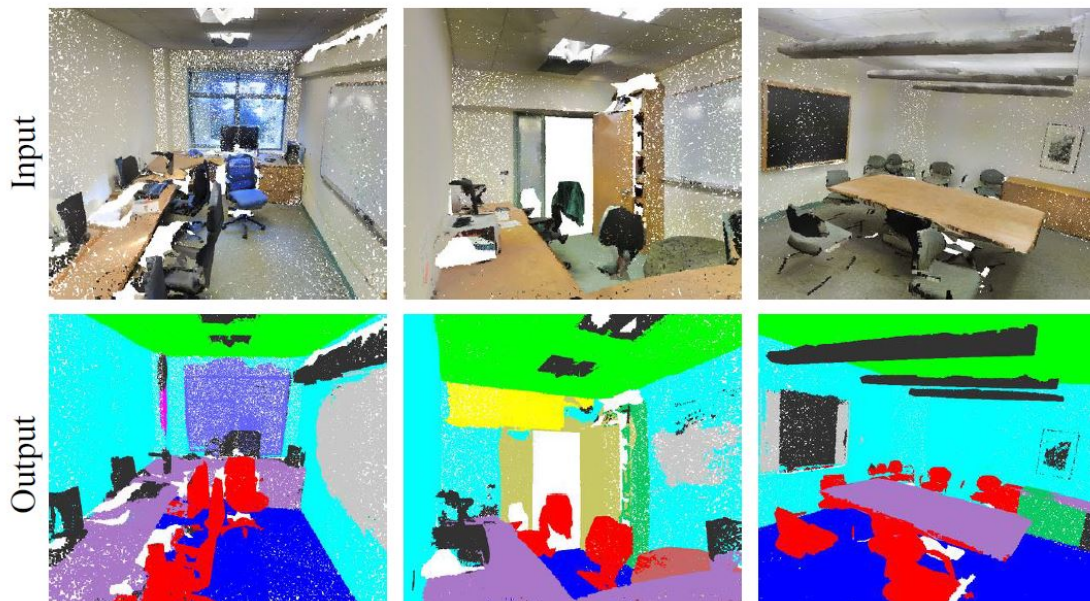


- Extract local features - Describes each point separately
- Extract global feature - Describes the entire point cloud
- Concatenate the local and global features and feed it into a shared MLP - The MLP learns to process the point feature according to a condition. The condition is described by the global feature vector.

Code (<https://github.com/nikitakaraevv/pointnet/blob/master/nbs/PointNetSeg.ipynb>) - Full implementation of PointNet Segmentation (can be opened in Colab)

Qi, Charles R., et al. "PointNet: Deep learning on point sets for 3d classification and segmentation." CVPR 2017.

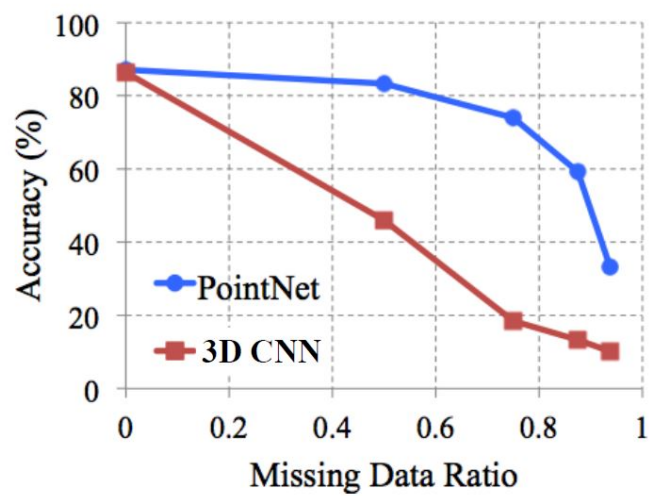
## Semantic Scene Parsing



Qi, Charles R., et al. "PointNet: Deep learning on point sets for 3d classification and segmentation." CVPR 2017.

### Results - Robustness to Missing Data (Classification example)

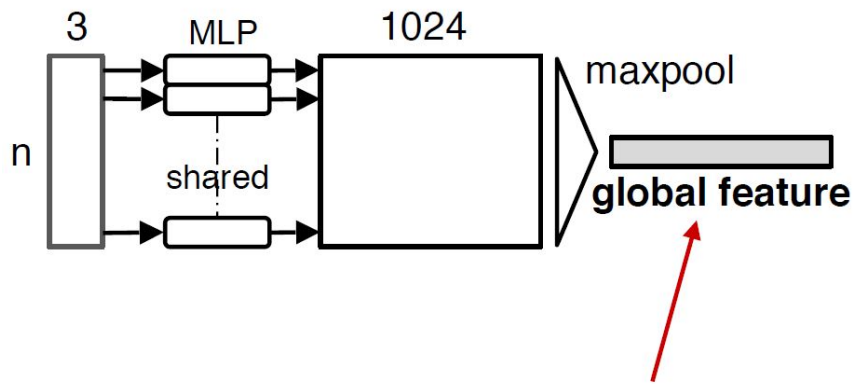
- Why is PointNet so robust to missing data?



Qi, Charles R., et al. "PointNet: Deep learning on point sets for 3d classification and segmentation." CVPR 2017.



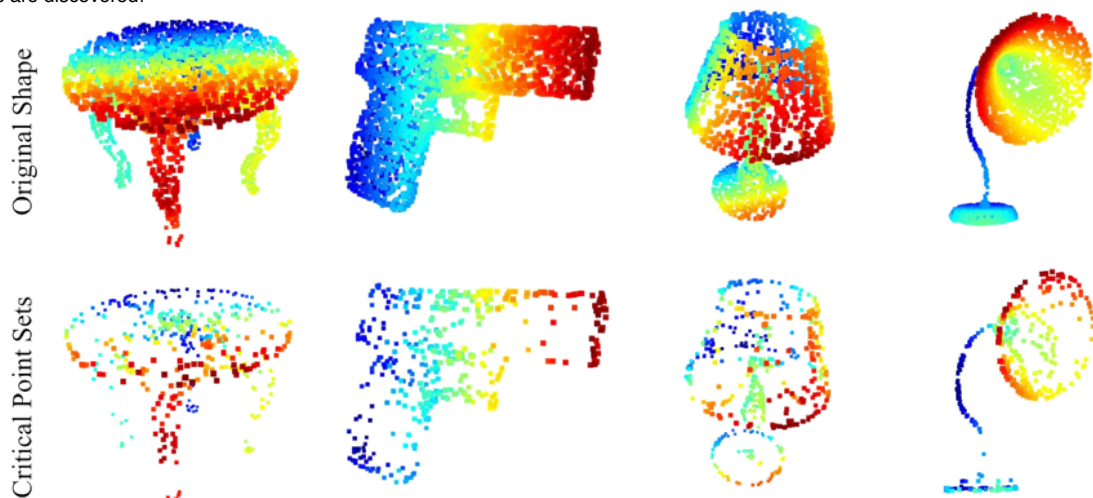
### Visualizing Global Point Cloud Features



Which input points are contributing to the global feature?  
**(critical points)**

### Visualize What is Learned by Reconstruction

- Salient points are discovered!

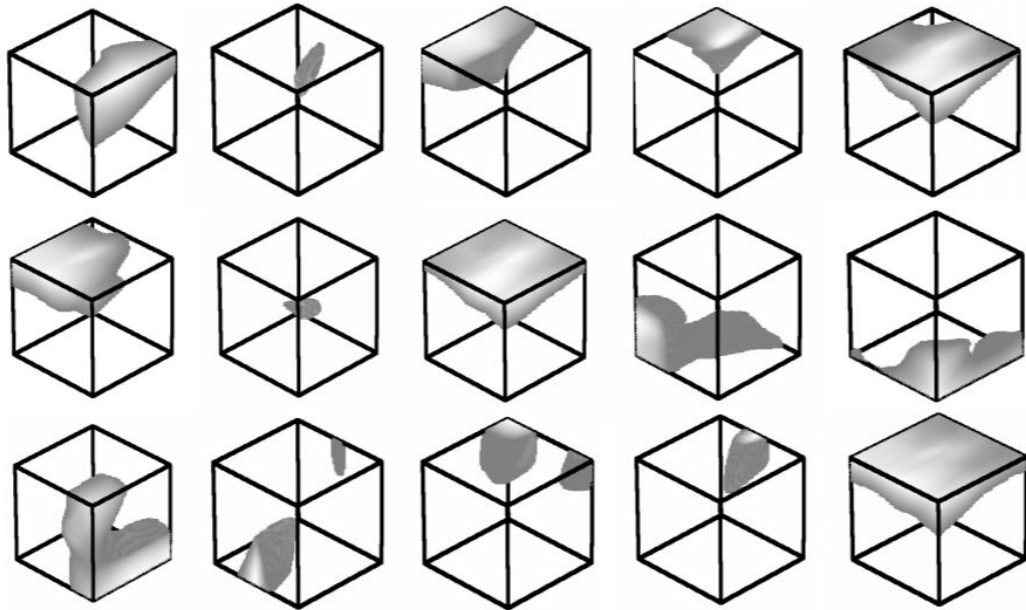


The "critical points" are those who influenced the global feature vector, a.k.a the pooling layer. The "critical" object's geometry structured is reserved.

### Point function visualization

For each per-point function  $h$  (MLP), calculate the values of  $h(p)$  for all the points  $p$  in the cube.

Random 15 function out of the 1024 learned functions:

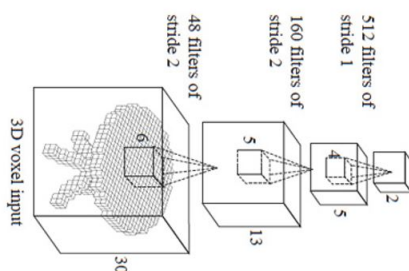


- Semi-equivalent to filter response in CNNs

Qi, Charles R., et al. "PointNet: Deep learning on point sets for 3d classification and segmentation." CVPR 2017.

### NA Limitations of PointNet

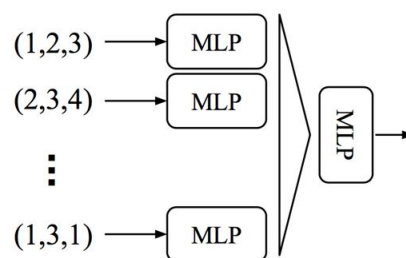
Hierarchical feature learning  
Multiple levels of abstraction



**3D CNN (Wu et al.)**

**v.s.**

Global feature learning  
Either one point or all points



**PointNet (vanilla) (Qi et al.)**

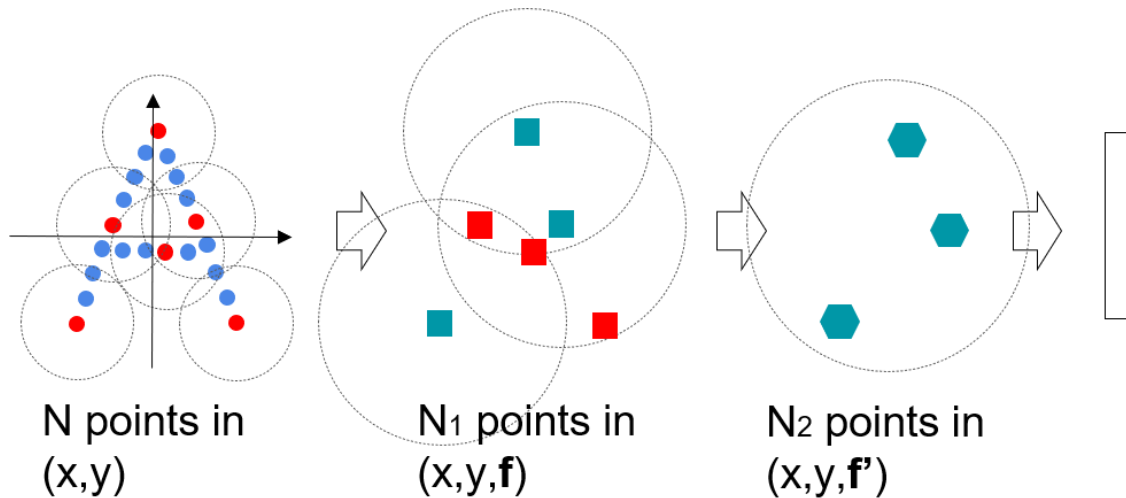
- No local context for each point
- Global feature depends on **absolute** coordinate. Hard to generalize to unseen scene configurations

### Points in Metric Space

- Learn “kernels” in 3D space and conduct convolution
- Kernels have compact spatial support
- For convolution, we need to find neighboring points
- Possible strategies for range query
  - Ball query (results in more stable features)
  - k-NN query (faster)



### PointNet v2.0: Multi-Scale PointNet



Repeated layers:

- Sample anchor points
- Find neighborhood of anchor points
- Apply PointNet in each neighborhood to mimic convolution

Qi, Charles Ruizhongtai, et al. "Pointnet++: Deep hierarchical feature learning on point sets in a metric space." Advances in neural information processing systems. 2017.

### More Point Clouds DL solutions:

- 3DmFV
- Dynamic Graph CNN
- PCNN
- PointCNN
- KPConv

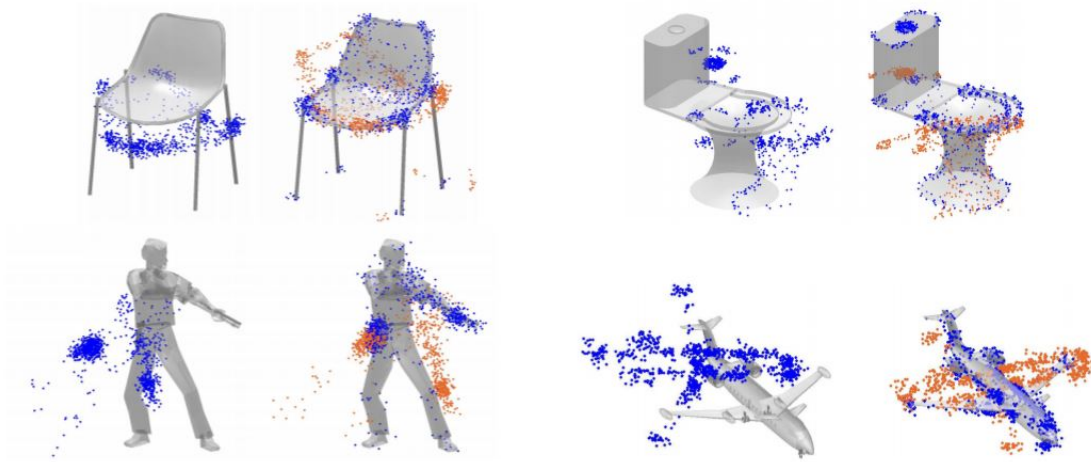


### 3D Deep Learning Applications

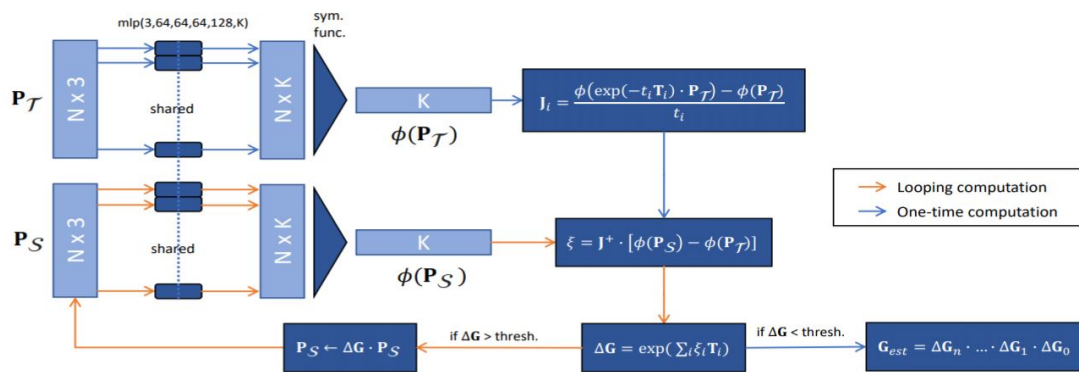
- Classification (V)
- Semantic segmentation (V)
- Part segmentation
- Object detection (Upcoming)
- Reconstruction
- Generation (Upcoming)
- Registration (Upcoming)
- Sampling - Downsampling, Upsampling
- SLAM
- Normal Estimation
- and many more...

## Registration:

Problem statment: Find the rotation and translation transformation between objects



- PointNetLK (blue) - Deep Learning, based on Lucas–Kanade method (Tracking lecture)
  - Comparing 2 point clouds using PointNet features
- ICP (orange) - Classic registration method

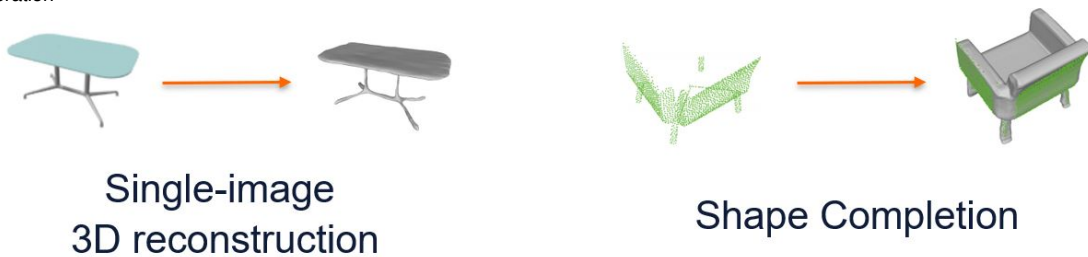


Both inputs (target and source) are being processed by PointNet architecture

Aoki, Yasuhiro, et al. "PointNetLK: Robust & efficient point cloud registration using PointNet." CVPR 2019.

## Generation

Conditional generation



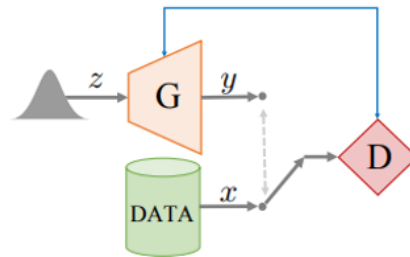
Free generation



*How would you build a point cloud GAN?*

**Learning Representations and Generative Models for 3D Point Clouds (Achlioptas et al.)**

- FC layer as generator
- PointNet as discriminator



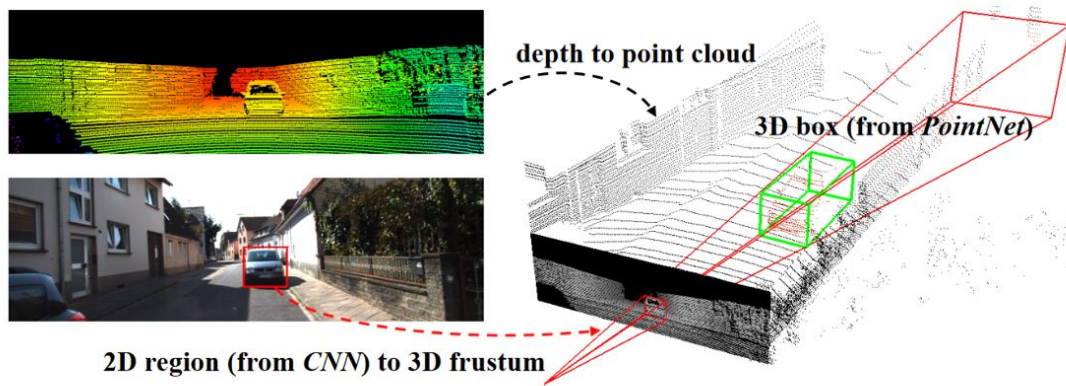
Achlioptas et al., "Learning Representations and Generative Models for 3D Point Clouds", ICML 2018

More generation methods:

- AtlasNet
- FoldingNet
- PointFlow
- OccupancyNetworks
- DeepSDF
- ...

### Detection:

- Generate object proposals from a view (e.g., using SSD)
- Recognize using PointNet



Qi et al., "Frustum PointNets for 3D Object Detection from RGB-D Data", CVPR 2018



### Questions

- What are the differences between 2D image and a point cloud?
  - Unstructured
  - Vary number of points
  - Unordered
- Why it might be hard to feed a point cloud as neural network (NN) input?
  - Does not rely on a grid
  - Does not have a fixed size
  - Different permutation represent the same point cloud

All three differences influence directly the ability of using NN!

- What are the benefits of using a point cloud?
  - Most sensors raw outputs are point clouds (LiDAR)
  - Very efficient representation of 3D data (no empty voxels)
  - Reserve geometric details (no quantization)



## Recommended Tools

---

Python:

- Open3D
- trimesh
- Ipyvolume - Visualization for Notebooks

Deep Learning:

- Python3D
- Kaolin (Pytorch)
- TensorFlow Graphics

Visualize Tools (drop and view):

- CloudCompare
- MeshLab

For more 3D deep learning frameworks and datasets:

- [awesome-point-cloud-analysis](https://github.com/Yochengliu/awesome-point-cloud-analysis) (<https://github.com/Yochengliu/awesome-point-cloud-analysis>).
- [3D-Machine-Learning#datasets](https://github.com/timzhang642/3D-Machine-Learning#datasets) (<https://github.com/timzhang642/3D-Machine-Learning#datasets>).

Datasets:

- ModelNet
- ShapeNet
- PartNet
- Sydney Urban Object Dataset
- Stanford 3D
- KITTI
- ...



## Recommended Videos

---



### Warning!

- These videos do not replace the lectures and tutorials.
- Please use these to get a better understanding of the material, and not as an alternative to the written material.

### Video By Subject

- 3D Deep Learning
  - General (Both highly recommended):
    - [3D Deep Learning Tutorial from SU lab at UCSD](https://www.youtube.com/watch?time_continue=6&v=vfL6uJYFrp4&feature=emb_logo) ([https://www.youtube.com/watch?time\\_continue=6&v=vfL6uJYFrp4&feature=emb\\_logo](https://www.youtube.com/watch?time_continue=6&v=vfL6uJYFrp4&feature=emb_logo)) - Hao Su
    - [Geometric deep learning](https://www.youtube.com/watch?v=wLU4YsC_4NYo) ([https://www.youtube.com/watch?v=wLU4YsC\\_4NYo](https://www.youtube.com/watch?v=wLU4YsC_4NYo)) - Micahel Bronstein
  - [PointNet](https://www.youtube.com/watch?v=Cge-hot0Oc0&t=24s) (<https://www.youtube.com/watch?v=Cge-hot0Oc0&t=24s>)
  - [3DmFV](https://www.youtube.com/watch?v=HIUGOKSLTcE) (<https://www.youtube.com/watch?v=HIUGOKSLTcE>)



## Credits

---

- Slides - [Yizhak \(Itzik\) Ben-Shabat](http://www.itzikbs.com/category/research-blog/) (<http://www.itzikbs.com/category/research-blog/>), [Simon Lucey \(CMU\)](https://ci2cv.net/people/simon-lucey/) (<https://ci2cv.net/people/simon-lucey/>), [Hao Su, Jiayuan Gu and Minghua Liu \(UCSanDiego\)](https://cseweb.ucsd.edu/~haosu/) (<https://cseweb.ucsd.edu/~haosu/>)
- Multiple View Geometry in Computer Vision - Hartley and Zisserman - Sections 9,10
- [Computer Vision: Algorithms and Applications](https://www.springer.com/gp/book/9781848829343) (<https://www.springer.com/gp/book/9781848829343>) - Richard Szeliski - Sections 11,12
- Icons from [Icon8.com](https://icons8.com/) (<https://icons8.com/>), - <https://icons8.com> (<https://icons8.com>)